Random Forest Regressor: Predicting Football Player Market Values

Mehmet Yağız Ceyhan ID:123200093

Date: 03/25/2026

Made with GAMMA

Concept Overview: The Wisdom of the Forest

How Random Forest Predicts Market Value?

Core Intuition: "Wisdom of the Crowd"

Instead of relying on one giant, complex tree (which overfits), we build a "forest" of many simple trees.

Made with GAMMA

The Process

- **Bootstrap:** Create random subsets of player data.
- **Train:** Build a separate Decision Tree for each subset.
- **Aggregate:** Average all predictions to find the final Market Value.

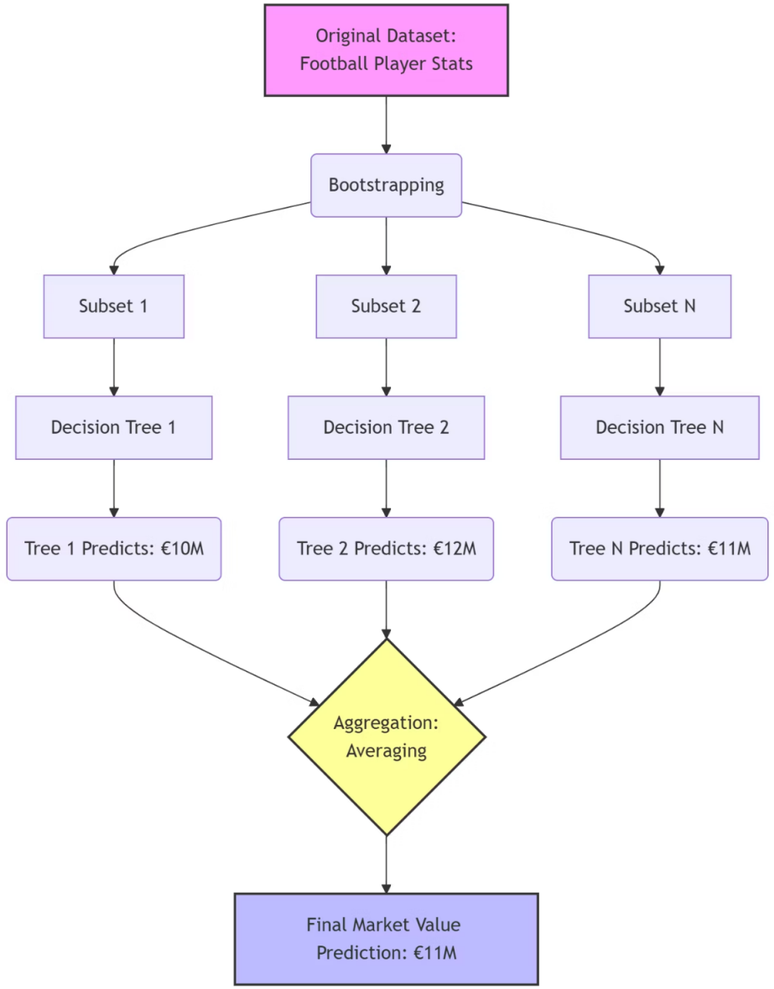

Made with GAMMA

"The Math: Splitting & Aggregation"

Minimizing Error & Maximizing Accuracy

- Splitting
- Ensemble

$$MSE = \frac{1}{n} \sum_{i=1}^{n} (Y_i - \hat{Y}_i)^2$$

Objective: Minimize node impurity

$$\hat{Y}_{rf} = \frac{1}{B} \sum_{b=1}^{B} f_b(x)$$

Objective: Average the forest's predictions

Made with GAMMA

Training Process

The Power of "Bagging" (Bootstrap Aggregation)

- **Bootstrapping**: The model creates multiple subsets from the original dataset by sampling *with replacement*.
- **Feature Randomness (Splitting)**: At each node of a tree, it doesn't look at all features (Age, Potential, Pace...). It randomly selects a small subset of features to find the best split.
- **Independent Growth (Parallel Training):** Unlike Boosting, trees are built independently and in parallel. They don't learn from each other's mistakes.

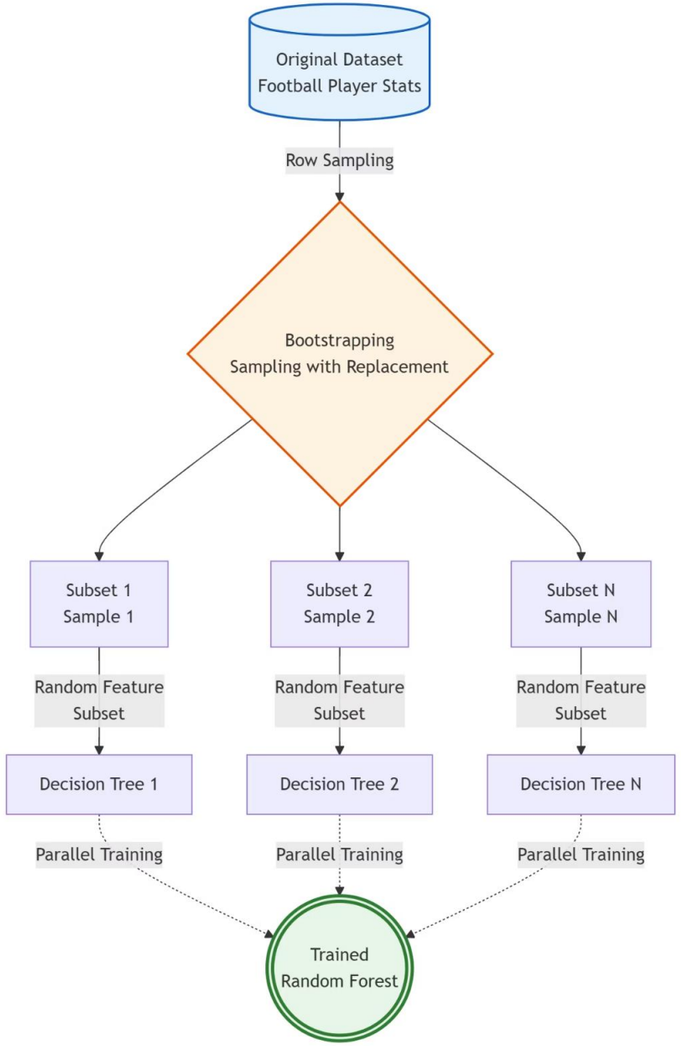

Made with GAMMA

Inference Mechanism

Predicting New Player Values

- **The Input:**A new, unseen player's data (e.g., Age 21, Potential 88, Goals 15) is fed into the trained forest.
- **Independent Predictions**: Each decision tree in the forest independently evaluates the player's stats and outputs its own continuous prediction (e.g., Tree 1 says €45M, Tree 2 says €50M).
- **Aggregation**: For regression tasks, the forest calculates the **mean average** of all individual tree predictions to output the final estimated market value.

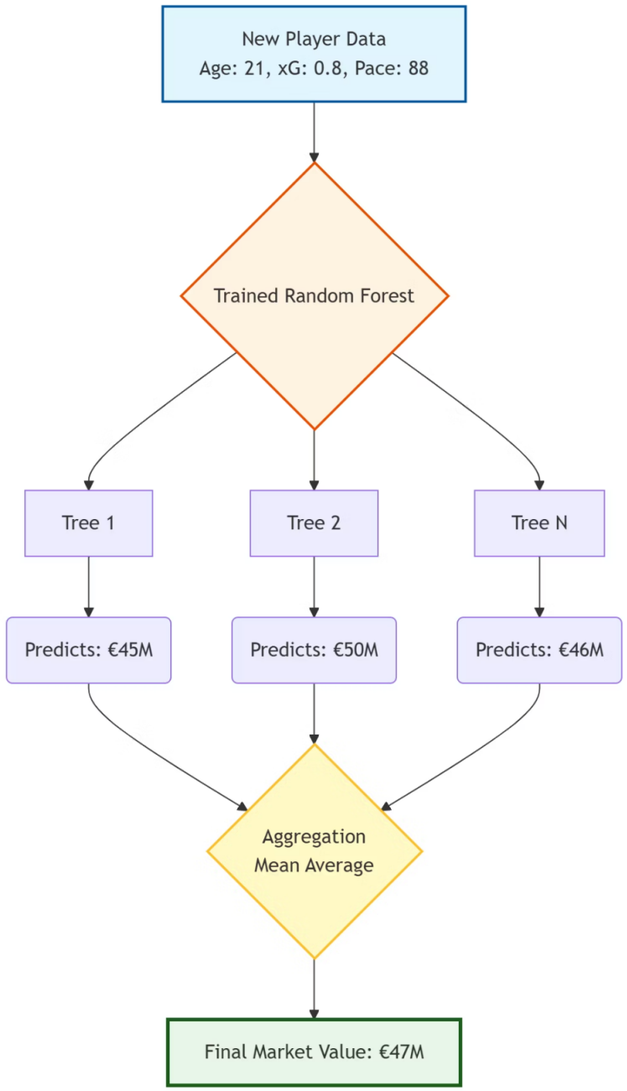

Made with GAMMA

Illustration Example

Feature Splitting in Action

- **Root Node**: The tree starts with the most critical feature that minimizes the MSE the most (e.g., *Potential*).
- **Internal Nodes:** The data is further split using secondary features like *Age* or *Overall Rating*.
- **Leaf Nodes:** The final stopping point contains the predicted **Market Value** for players falling into that specific path.

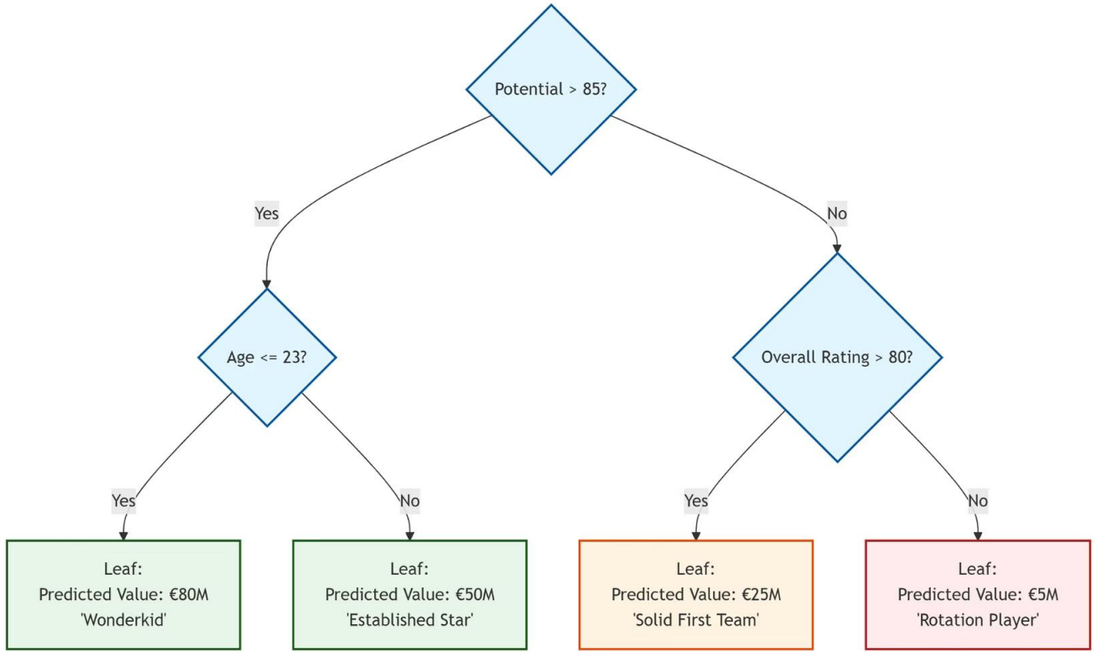

Made with GAMMA

Pros

- **Robust Against Overfitting**: By averaging multiple trees (bagging), it prevents the model from memorizing the training data.
- **Feature Importance Built-In:** Automatically calculates which features (e.g., Potential, Age) impact the Market Value the most
- **Handles Non-Linearity Well**: Easily captures complex relationships in player stats without needing extensive feature engineering.

Cons:

- **Computationally Expensive**: Training hundreds of deep trees can be slow and requires significant memory.
- **"Black Box" Nature:** Harder to interpret the exact path of a prediction compared to a single, simple Decision Tree or Logistic Regression
- **Poor at Extrapolation**: It cannot predict values outside the range of the training data (e.g., predicting a €500M transfer if the max in data was €200M).

Made with GAMMA